

Tight Confidence Sequences From Universal Portfolio

Based on Orabona and Jun (2024), IEEE Transactions on Information Theory

Xiaoqi (Shirley) Liu

E-Values Reading Group
Department of Statistics, University of Oxford

10 March 2026

Problem setup

Problem setup

We observe bounded data $X_1, X_2, \dots$ with

$$0 \leq X_t \leq 1, \quad \mathbb{E}[X_t \mid X_{1:t-1}] = \mu \quad \text{a.s.}$$

Problem setup

We observe bounded data $X_1, X_2, \dots$ with

$$0 \leq X_t \leq 1, \quad \mathbb{E}[X_t \mid X_{1:t-1}] = \mu \quad \text{a.s.}$$

Goal: a confidence sequence (CS), i.e., a sequence of confidence intervals $I_t = [\ell_t(X_{1:t}), u_t(X_{1:t})]$ such that

$$\mathbb{P}(\mu \in I_t, \forall t \geq 1) \geq 1 - \delta.$$

Problem setup

We observe bounded data $X_1, X_2, \dots$ with

$$0 \leq X_t \leq 1, \quad \mathbb{E}[X_t \mid X_{1:t-1}] = \mu \quad \text{a.s.}$$

Goal: a confidence sequence (CS), i.e., a sequence of confidence intervals $I_t = [\ell_t(X_{1:t}), u_t(X_{1:t})]$ such that

$$\mathbb{P}(\mu \in I_t, \forall t \geq 1) \geq 1 - \delta.$$

Main contribution of the paper: tight CSs based on regret bounds for **betting/universal portfolio** algorithms.

Outline

1. Recap and warm-up

- 1.1 Testing by betting
- 1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

- 2.1 A reformulation using online learning tools
- 2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

Outline

1. Recap and warm-up

1.1 Testing by betting

1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

2.1 A reformulation using online learning tools

2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

Testing by betting

Testing by betting

- Start with $\text{Wealth}_0 = 1$

Testing by betting

- Start with $\text{Wealth}_0 = 1$
- Bet on **continuous coin** outcomes: $c_t \in [-1, 1]$ (possibly adversarial)

Testing by betting

- Start with $\text{Wealth}_0 = 1$
- Bet on **continuous coin** outcomes: $c_t \in [-1, 1]$ (possibly adversarial)
- Bet a signed fraction $\beta_t \in [-1, 1]$ of current wealth on coin being positive ($\beta_t > 0$) or negative ($\beta_t < 0$). Update wealth

$$\text{Wealth}_t = \text{Wealth}_{t-1}(1 + \beta_t c_t).$$

Testing by betting

- Start with $\text{Wealth}_0 = 1$
- Bet on **continuous coin** outcomes: $c_t \in [-1, 1]$ (possibly adversarial)
- Bet a signed fraction $\beta_t \in [-1, 1]$ of current wealth on coin being positive ($\beta_t > 0$) or negative ($\beta_t < 0$). Update wealth

$$\text{Wealth}_t = \text{Wealth}_{t-1}(1 + \beta_t c_t).$$

Under H_0 coin is fair, i.e., $\mathbb{E}[c_t \mid c_{1:t-1}] = 0$, and $(\text{Wealth}_i)_{i=1}^t$ is a nonnegative supermartingale, i.e.,

$$\mathbb{E}[\text{Wealth}_t \mid \text{Wealth}_{1:t-1}] = \text{Wealth}_{t-1}.$$

Testing by betting

- Start with $\text{Wealth}_0 = 1$
- Bet on **continuous coin** outcomes: $c_t \in [-1, 1]$ (possibly adversarial)
- Bet a signed fraction $\beta_t \in [-1, 1]$ of current wealth on coin being positive ($\beta_t > 0$) or negative ($\beta_t < 0$). Update wealth

$$\text{Wealth}_t = \text{Wealth}_{t-1}(1 + \beta_t c_t).$$

Under H_0 coin is fair, i.e., $\mathbb{E}[c_t \mid c_{1:t-1}] = 0$, and $(\text{Wealth}_i)_{i=1}^t$ is a nonnegative supermartingale, i.e.,

$$\mathbb{E}[\text{Wealth}_t \mid \text{Wealth}_{1:t-1}] = \text{Wealth}_{t-1}.$$

Thus by **Ville's inequality**, we obtain the **time-uniform** error guarantee:

$$\mathbb{P}\left(\sup_{t \geq 0} \text{Wealth}_t \geq \frac{1}{\delta}\right) \leq \delta.$$

Testing by betting

- Start with $\text{Wealth}_0 = 1$
- Bet on **continuous coin** outcomes: $c_t \in [-1, 1]$ (possibly adversarial)
- Bet a signed fraction $\beta_t \in [-1, 1]$ of current wealth on coin being positive ($\beta_t > 0$) or negative ($\beta_t < 0$). Update wealth

$$\text{Wealth}_t = \text{Wealth}_{t-1}(1 + \beta_t c_t).$$

Under H_0 coin is fair, i.e., $\mathbb{E}[c_t \mid c_{1:t-1}] = 0$, and $(\text{Wealth}_i)_{i=1}^t$ is a nonnegative supermartingale, i.e.,

$$\mathbb{E}[\text{Wealth}_t \mid \text{Wealth}_{1:t-1}] = \text{Wealth}_{t-1}.$$

Thus by **Ville's inequality**, we obtain the **time-uniform** error guarantee:

$$\mathbb{P}\left(\sup_{t \geq 0} \text{Wealth}_t \geq \frac{1}{\delta}\right) \leq \delta.$$

Testing by betting can be generalised through e-values & e-processes (Ramdas and Wang, 2025)

From testing by betting to mean inference

From testing by betting to mean inference

Consider random variables $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_t \mid X_{1:t-1}] = \mu$ a.s..

From testing by betting to mean inference

Consider random variables $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_t \mid X_{1:t-1}] = \mu$ a.s..

For each **candidate mean** $m \in [0, 1]$, define the continuous coin

$$c_t(m) := X_t - m \in [-m, 1 - m] \subseteq [-1, 1].$$

From testing by betting to mean inference

Consider random variables $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_t \mid X_{1:t-1}] = \mu$ a.s..

For each **candidate mean** $m \in [0, 1]$, define the continuous coin

$$c_t(m) := X_t - m \in [-m, 1 - m] \subseteq [-1, 1].$$

Under H_0 , $\mu = m$, i.e., $\mathbb{E}[c_t(m) \mid X_{1:t-1}] = 0$.

From testing by betting to mean inference

Consider random variables $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_t \mid X_{1:t-1}] = \mu$ a.s..

For each **candidate mean** $m \in [0, 1]$, define the continuous coin

$$c_t(m) := X_t - m \in [-m, 1 - m] \subseteq [-1, 1].$$

Under H_0 , $\mu = m$, i.e., $\mathbb{E}[c_t(m) \mid X_{1:t-1}] = 0$.

Build a wealth process $W_t(m)$ for each m , where by Ville,

$$\mathbb{P}\left(\exists t \geq 0 : \text{Wealth}_t(\mu) \geq \frac{1}{\delta}\right) \leq \delta.$$

From testing by betting to mean inference

Consider random variables $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_t \mid X_{1:t-1}] = \mu$ a.s..

For each **candidate mean** $m \in [0, 1]$, define the continuous coin

$$c_t(m) := X_t - m \in [-m, 1 - m] \subseteq [-1, 1].$$

Under H_0 , $\mu = m$, i.e., $\mathbb{E}[c_t(m) \mid X_{1:t-1}] = 0$.

Build a wealth process $W_t(m)$ for each m , where by Ville,

$$\mathbb{P}\left(\exists t \geq 0 : \text{Wealth}_t(\mu) \geq \frac{1}{\delta}\right) \leq \delta.$$

Define the CS as the set of $m \in [0, 1]$ that **haven't been rejected yet**:

$$I_t := \left\{ m \in [0, 1] : \sup_{0 \leq s \leq t} \text{Wealth}_s(m) < \frac{1}{\delta} \right\}.$$

From testing by betting to mean inference

Consider random variables $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_t \mid X_{1:t-1}] = \mu$ a.s..

For each **candidate mean** $m \in [0, 1]$, define the continuous coin

$$c_t(m) := X_t - m \in [-m, 1 - m] \subseteq [-1, 1].$$

Under H_0 , $\mu = m$, i.e., $\mathbb{E}[c_t(m) \mid X_{1:t-1}] = 0$.

Build a wealth process $W_t(m)$ for each m , where by Ville,

$$\mathbb{P}\left(\exists t \geq 0 : \text{Wealth}_t(\mu) \geq \frac{1}{\delta}\right) \leq \delta.$$

Define the CS as the set of $m \in [0, 1]$ that **haven't been rejected yet**:

$$I_t := \left\{ m \in [0, 1] : \sup_{0 \leq s \leq t} \text{Wealth}_s(m) < \frac{1}{\delta} \right\}.$$

⇒ Which betting algorithm should we use?

Outline

1. Recap and warm-up

1.1 Testing by betting

1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

2.1 A reformulation using online learning tools

2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

Warm-up: the Krichevsky–Trofimov (KT) betting algorithm

Warm-up: the Krichevsky–Trofimov (KT) betting algorithm

A classic choice is the **KT bettor** that bets

$$\beta_t = \frac{\sum_{i=1}^{t-1} c_i}{t}.$$

Warm-up: the Krichevsky–Trofimov (KT) betting algorithm

A classic choice is the **KT bettor** that bets

$$\beta_t = \frac{\sum_{i=1}^{t-1} c_i}{t}.$$

Worst-case guarantee (Orabona and Pal, 2016): for *any* sequence of $c_1, c_2, \dots \in [-1, 1]$, we have

$$\ln \text{Wealth}_t \geq \frac{(\sum_{i=1}^t c_i)^2}{2t} - \frac{1}{2} \ln(4t).$$

Warm-up: the Krichevsky–Trofimov (KT) betting algorithm

A classic choice is the **KT bettor** that bets

$$\beta_t = \frac{\sum_{i=1}^{t-1} c_i}{t}.$$

Worst-case guarantee (Orabona and Pal, 2016): for *any* sequence of $c_1, c_2, \dots \in [-1, 1]$, we have

$$\ln \text{Wealth}_t \geq \frac{(\sum_{i=1}^t c_i)^2}{2t} - \frac{1}{2} \ln(4t).$$

Plug in $c_t = X_t - \mu$ and combine with Ville:

$$\mathbb{P}\left(\exists t \geq 0 : \frac{(\sum_{i=1}^t (X_i - \mu))^2}{2t} \geq \ln\left(\frac{1}{\delta}\right) + \frac{1}{2} \ln(4t)\right) \leq \delta.$$

Warm-up: the Krichevsky–Trofimov (KT) betting algorithm

A classic choice is the **KT bettor** that bets

$$\beta_t = \frac{\sum_{i=1}^{t-1} c_i}{t}.$$

Worst-case guarantee (Orabona and Pal, 2016): for *any* sequence of $c_1, c_2, \dots \in [-1, 1]$, we have

$$\ln \text{Wealth}_t \geq \frac{(\sum_{i=1}^t c_i)^2}{2t} - \frac{1}{2} \ln(4t).$$

Plug in $c_t = X_t - \mu$ and combine with Ville:

$$\mathbb{P}\left(\exists t \geq 0 : \frac{(\sum_{i=1}^t (X_i - \mu))^2}{2t} \geq \ln\left(\frac{1}{\delta}\right) + \frac{1}{2} \ln(4t)\right) \leq \delta.$$

Inverting the inequality gives

$$\mathbb{P}\left(\exists t \geq 0 : \left|\mu - \frac{1}{t} \sum_{i=1}^t X_i\right| \geq \sqrt{\frac{2 \ln \frac{2\sqrt{t}}{\delta}}{t}}\right) \leq \delta.$$

Why KT is not the end of the story

Why KT is not the end of the story

Issues with KT when we have $c_t(m) = X_t - m \in [-m, 1 - m] \subseteq [-1, 1]$:

Why KT is not the end of the story

Issues with KT when we have $c_t(m) = X_t - m \in [-m, 1 - m] \subseteq [-1, 1]$:

- KT gives a time-uniform **Hoeffding-type** bound. It only uses boundedness, ignoring **variance**.

Why KT is not the end of the story

Issues with KT when we have $c_t(m) = X_t - m \in [-m, 1 - m] \subseteq [-1, 1]$:

- KT gives a time-uniform **Hoeffding-type** bound. It only uses boundedness, ignoring **variance**.
- KT bets lie in the small interval $\beta_t = \frac{1}{t} \sum_{i=1}^{t-1} c_i \in [-1, 1]$, whereas the admissible betting range is $\beta_t \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$ due to $1 + \beta_t c_t(m) \geq 0$ for all $c_t(m) \in [-m, 1 - m]$.

Why KT is not the end of the story


Issues with KT when we have $c_t(m) = X_t - m \in [-m, 1 - m] \subseteq [-1, 1]$:

- KT gives a time-uniform **Hoeffding-type** bound. It only uses boundedness, ignoring **variance**.
- KT bets lie in the small interval $\beta_t = \frac{1}{t} \sum_{i=1}^{t-1} c_i \in [-1, 1]$, whereas the admissible betting range is $\beta_t \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$ due to $1 + \beta_t c_t(m) \geq 0$ for all $c_t(m) \in [-m, 1 - m]$.
 $\Rightarrow$ KT **under-utilises** the allowable range near the boundaries $m = -1$ or 1 .

Why KT is not the end of the story

Issues with KT when we have $c_t(m) = X_t - m \in [-m, 1 - m] \subseteq [-1, 1]$:

- KT gives a time-uniform **Hoeffding-type** bound. It only uses boundedness, ignoring **variance**.
 - KT bets lie in the small interval $\beta_t = \frac{1}{t} \sum_{i=1}^{t-1} c_i \in [-1, 1]$, whereas the admissible betting range is $\beta_t \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$ due to $1 + \beta_t c_t(m) \geq 0$ for all $c_t(m) \in [-m, 1 - m]$.
- $\Rightarrow$ KT **under-utilises** the allowable range near the boundaries $m = 1$ or 0 .




Why KT is not the end of the story

Issues with KT when we have $c_t(m) = X_t - m \in [-m, 1 - m] \subseteq [-1, 1]$:

- KT gives a time-uniform **Hoeffding-type** bound. It only uses boundedness, ignoring **variance**.
- KT bets lie in the small interval $\beta_t = \frac{1}{t} \sum_{i=1}^{t-1} c_i \in [-1, 1]$, whereas the admissible betting range is $\beta_t \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$ due to $1 + \beta_t c_t(m) \geq 0$ for all $c_t(m) \in [-m, 1 - m]$.
 $\Rightarrow$ KT **under-utilises** the allowable range near the boundaries $m = 1$ or 0 .

$\Rightarrow$ **Can we do better than KT?**



Outline

1. Recap and warm-up

- 1.1 Testing by betting
- 1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

- 2.1 A reformulation using online learning tools
- 2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

Outline

1. Recap and warm-up

- 1.1 Testing by betting
- 1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

- 2.1 A reformulation using online learning tools
- 2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

A reformulation using online learning tools

A reformulation using online learning tools

To build CS, for each $t \geq 0$, collect $m \in [0, 1]$ such that

A reformulation using online learning tools

To build CS, for each $t \geq 0$, collect $m \in [0, 1]$ such that

$$\ln \text{Wealth}_t(\beta, m) < \ln \frac{1}{\delta}.$$

A reformulation using online learning tools

To build CS, for each $t \geq 0$, collect $m \in [0, 1]$ such that

$$\ln \text{Wealth}_t(\beta_t^*, m) - \underbrace{\left(\ln \text{Wealth}_t(\beta_t^*, m) - \ln \text{Wealth}_t(\beta, m) \right)}_{=:\text{Regret}_t(m) \text{ (associated with log-wealth)}} < \ln \frac{1}{\delta},$$

where β_t^* is best fixed bet in hindsight:

$$\beta_t^*(m) \in \arg \max_{\beta \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]} \ln \text{Wealth}_t(\beta, m), \quad \ln \text{Wealth}_t(\beta, m) = \sum_{i=1}^t \ln(1 + \beta(X_i - m)).$$

A reformulation using online learning tools

To build CS, for each $t \geq 0$, collect $m \in [0, 1]$ such that

$$\ln \text{Wealth}_t(\beta_t^*, m) - \underbrace{\left(\ln \text{Wealth}_t(\beta_t^*, m) - \ln \text{Wealth}_t(\beta, m) \right)}_{=:\text{Regret}_t(m) \text{ (associated with log-wealth)}} < \ln \frac{1}{\delta},$$

where β_t^* is best fixed bet in hindsight:

$$\beta_t^*(m) \in \arg \max_{\beta \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]} \ln \text{Wealth}_t(\beta, m), \quad \ln \text{Wealth}_t(\beta, m) = \sum_{i=1}^t \ln(1 + \beta(X_i - m)).$$

The paper proves inequalities of the schematic form

$$\mathbb{P}\left(\ln \text{Wealth}_t(\beta_t^*, \mu) \leq \text{Regret}_t(\mu) + \ln \frac{1}{\delta}, \forall t \geq 0 \right) \geq 1 - \delta.$$

A reformulation using online learning tools

To build CS, for each $t \geq 0$, collect $m \in [0, 1]$ such that

$$\ln \text{Wealth}_t(\beta_t^*, m) - \underbrace{\left(\ln \text{Wealth}_t(\beta_t^*, m) - \ln \text{Wealth}_t(\beta, m) \right)}_{=:\text{Regret}_t(m) \text{ (associated with log-wealth)}} < \ln \frac{1}{\delta},$$

where β_t^* is best fixed bet in hindsight:

$$\beta_t^*(m) \in \arg \max_{\beta \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]} \ln \text{Wealth}_t(\beta, m), \quad \ln \text{Wealth}_t(\beta, m) = \sum_{i=1}^t \ln(1 + \beta(X_i - m)).$$

The paper proves inequalities of the schematic form

$$\mathbb{P}\left(\ln \text{Wealth}_t(\beta_t^*, \mu) \leq \text{Regret}_t(\mu) + \ln \frac{1}{\delta}, \forall t \geq 0 \right) \geq 1 - \delta.$$

Issue: near-optimal betting algorithms are often computationally infeasible.

A reformulation using online learning tools

To build CS, for each $t \geq 0$, collect $m \in [0, 1]$ such that

$$\ln \text{Wealth}_t(\beta_t^*, m) - \underbrace{\left(\ln \text{Wealth}_t(\beta_t^*, m) - \ln \text{Wealth}_t(\beta, m) \right)}_{=:\text{Regret}_t(m) \text{ (associated with log-wealth)}} < \ln \frac{1}{\delta},$$

where β_t^* is best fixed bet in hindsight:

$$\beta_t^*(m) \in \arg \max_{\beta \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]} \ln \text{Wealth}_t(\beta, m), \quad \ln \text{Wealth}_t(\beta, m) = \sum_{i=1}^t \ln(1 + \beta(X_i - m)).$$

The paper proves inequalities of the schematic form

$$\mathbb{P}\left(\ln \text{Wealth}_t(\beta_t^*, \mu) \leq \text{Regret}_t(\mu) + \ln \frac{1}{\delta}, \forall t \geq 0 \right) \geq 1 - \delta.$$

Issue: near-optimal betting algorithms are often computationally infeasible.

Remedy: Instead of regret, plug in a **regret upper bound** $R_t(m)$, so no need to run the algorithm!

A reformulation using online learning tools

To build CS, for each $t \geq 0$, collect $m \in [0, 1]$ such that

$$\ln \text{Wealth}_t(\beta_t^*, m) - \underbrace{\left(\ln \text{Wealth}_t(\beta_t^*, m) - \ln \text{Wealth}_t(\beta, m) \right)}_{=:\text{Regret}_t(m) \text{ (associated with log-wealth)}} < \ln \frac{1}{\delta},$$

where β_t^* is best fixed bet in hindsight:

$$\beta_t^*(m) \in \arg \max_{\beta \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]} \ln \text{Wealth}_t(\beta, m), \quad \ln \text{Wealth}_t(\beta, m) = \sum_{i=1}^t \ln(1 + \beta(X_i - m)).$$

The paper proves inequalities of the schematic form

$$\mathbb{P}\left(\ln \text{Wealth}_t(\beta_t^*, \mu) \leq R_t(\mu) + \ln \frac{1}{\delta}, \forall t \geq 0\right) \geq 1 - \delta.$$

Issue: near-optimal betting algorithms are often computationally infeasible.

Remedy: Instead of regret, plug in a **regret upper bound** $R_t(m)$, so no need to run the algorithm!

Take-away message: Any betting algorithm with a guaranteed regret (whether or not computationally feasible) yields a CS.

Take-away message: Any betting algorithm with a guaranteed regret (whether or not computationally feasible) yields a CS.

→ Paper focuses on betting algorithms via a **reduction** to **online portfolio selection**.

Outline

1. Recap and warm-up

1.1 Testing by betting

1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

2.1 A reformulation using online learning tools

2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

From coin betting to 2-stock online portfolio selection

From coin betting to 2-stock online portfolio selection

Lemma 1. Let $m \in [0, 1]$ and $c \in [-m, 1 - m]$. Set two market gains $w_1 = 1 + \frac{c}{m}$ and $w_2 = 1 - \frac{c}{1-m}$ (note $w_1, w_2 \geq 0$ so they are valid market gains). Define $\mathbf{b} = [b, 1 - b]^\top$ to be the play of a 2-stocks portfolio algorithm, where $b \in [0, 1]$. Then, by taking

$$\beta = -\frac{1}{1-m} + \left(\frac{1}{1-m} + \frac{1}{m}\right)b \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$$

as the signed betting fraction, a continuous coin-betting algorithm on c ensures that the gain in the coin-betting problem **coincides with** the gain in the portfolio selection problem.

From coin betting to 2-stock online portfolio selection

Lemma 1. Let $m \in [0, 1]$ and $c \in [-m, 1 - m]$. Set two market gains $w_1 = 1 + \frac{c}{m}$ and $w_2 = 1 - \frac{c}{1-m}$ (note $w_1, w_2 \geq 0$ so they are valid market gains). Define $\mathbf{b} = [b, 1 - b]^\top$ to be the play of a 2-stocks portfolio algorithm, where $b \in [0, 1]$. Then, by taking

$$\beta = -\frac{1}{1-m} + \left(\frac{1}{1-m} + \frac{1}{m}\right)b \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$$

as the signed betting fraction, a continuous coin-betting algorithm on c ensures that the gain in the coin-betting problem **coincides with** the gain in the portfolio selection problem.

Proof sketch.

$$\begin{bmatrix} b \\ 1 - b \end{bmatrix} \cdot \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = b\left(1 + \frac{c}{m}\right) + (1 - b)\left(1 - \frac{c}{1-m}\right) = 1 + c\left(\frac{b}{m} - \frac{1-b}{1-m}\right) = 1 + c\beta.$$

From coin betting to 2-stock online portfolio selection

Lemma 1. Let $m \in [0, 1]$ and $c \in [-m, 1 - m]$. Set two market gains $w_1 = 1 + \frac{c}{m}$ and $w_2 = 1 - \frac{c}{1-m}$ (note $w_1, w_2 \geq 0$ so they are valid market gains). Define $\mathbf{b} = [b, 1 - b]^\top$ to be the play of a 2-stocks portfolio algorithm, where $b \in [0, 1]$. Then, by taking

$$\beta = -\frac{1}{1-m} + \left(\frac{1}{1-m} + \frac{1}{m}\right)b \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$$

as the signed betting fraction, a continuous coin-betting algorithm on c ensures that the gain in the coin-betting problem **coincides with** the gain in the portfolio selection problem.

Interpretation: Betting on an asymmetric bounded coin **is equivalent to** choosing a portfolio weight in a 2-stock online universal portfolio.

$$\Rightarrow \ln \text{Wealth}^{\text{coin}}(\beta) = \ln \text{Wealth}^{\text{portfolio}}(\mathbf{b}) \text{ and } \text{Regret}_t^{\text{coin}}(\beta) = \text{Regret}_t^{\text{portfolio}}(\mathbf{b})$$

Universal portfolio algorithms

Universal portfolio algorithms

A portfolio algorithm is **universal** if the regret against *any* sequence of market gain vectors $[w_1, w_2]^\top$ is sublinear in T (Cover and Ordentlich, 1996).

Universal portfolio algorithms

A portfolio algorithm is **universal** if the regret against *any* sequence of market gain vectors $[w_1, w_2]^\top$ is sublinear in T (Cover and Ordentlich, 1996).

Paper focuses on a mixture-based universal portfolio algorithm called the **F -weighted portfolio algorithm** (Cover, 1991)

Universal portfolio algorithms

A portfolio algorithm is **universal** if the regret against *any* sequence of market gain vectors $[w_1, w_2]^\top$ is sublinear in T (Cover and Ordentlich, 1996).

Paper focuses on a mixture-based universal portfolio algorithm called the **F -weighted portfolio algorithm** (Cover, 1991) that outputs at each step

$$\mathbf{b}_t = \frac{\int_B \mathbf{b} \text{Wealth}_{t-1}(\mathbf{b}) dF(\mathbf{b})}{\int_B \text{Wealth}_{t-1}(\mathbf{b}) dF(\mathbf{b})}, \quad B := \{[b, 1 - b]^\top \in \mathbb{R}^2 : b \in [0, 1]\},$$

Universal portfolio algorithms

A portfolio algorithm is **universal** if the regret against *any* sequence of market gain vectors $[w_1, w_2]^\top$ is sublinear in T (Cover and Ordentlich, 1996).

Paper focuses on a mixture-based universal portfolio algorithm called the **F -weighted portfolio algorithm** (Cover, 1991) that outputs at each step

$$\mathbf{b}_t = \frac{\int_B \mathbf{b} \text{Wealth}_{t-1}(\mathbf{b}) dF(\mathbf{b})}{\int_B \text{Wealth}_{t-1}(\mathbf{b}) dF(\mathbf{b})}, \quad B := \{[b, 1-b]^\top \in \mathbb{R}^2 : b \in [0, 1]\},$$

with regret guarantee

$$\text{Regret}_T = \max_{\mathbf{b} \in B} \ln \text{Wealth}_T(\mathbf{b}) - \ln \underbrace{\int_B \text{Wealth}_t(\mathbf{b}) dF(\mathbf{b})}_{\text{wealth of the } F\text{-weighted portfolio algorithm}} = o(T).$$

Outline

1. Recap and warm-up

1.1 Testing by betting

1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

2.1 A reformulation using online learning tools

2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

PRECiSE-CO96:

Portfolio REgret for Confidence SEquences Using Cover and Ordentlich (1996)

PRECiSE-CO96:

Portfolio REgret for Confidence SEquences Using Cover and Ordentlich (1996)

For the F -weighted universal portfolio algorithm with $F = \text{Beta}(\frac{1}{2}, \frac{1}{2})$, its regret is upper bounded by $\text{Regret}_t \leq \ln \frac{\sqrt{\pi}\Gamma(t+1)}{\Gamma(t+\frac{1}{2})} \leq \frac{1}{2} \ln(4t)$.

PRECiSE-CO96:

Portfolio REgret for Confidence SEquences Using Cover and Ordentlich (1996)

For the F -weighted universal portfolio algorithm with $F = \text{Beta}(\frac{1}{2}, \frac{1}{2})$, its regret is upper bounded by $\text{Regret}_t \leq \ln \frac{\sqrt{\pi}\Gamma(t+1)}{\Gamma(t+\frac{1}{2})} \leq \frac{1}{2} \ln(4t)$. This yields the following CS:

Theorem 1. Let $\delta \in (0, 1)$. Let $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_i \mid X_{1:i-1}] = \mu$ a.s.. Then, PRECiSE-CO96 yields

$$\mathbb{P}\left(\exists t \geq 1 : \max_{\beta \in \left[-\frac{1}{1-\mu}, \frac{1}{\mu}\right]} \sum_{i=1}^t \ln(1 + \beta(X_i - \mu)) \geq \ln\left(\frac{1}{\delta}\right) + \frac{1}{2} \ln(4t)\right) \leq \delta.$$

PRECiSE-CO96:

Portfolio REgret for Confidence SEquences Using Cover and Ordentlich (1996)

For the F -weighted universal portfolio algorithm with $F = \text{Beta}(\frac{1}{2}, \frac{1}{2})$, its regret is upper bounded by $\text{Regret}_t \leq \ln \frac{\sqrt{\pi} \Gamma(t+1)}{\Gamma(t+\frac{1}{2})} \leq \frac{1}{2} \ln(4t)$. This yields the following CS:

Theorem 1. Let $\delta \in (0, 1)$. Let $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_i \mid X_{1:i-1}] = \mu$ a.s.. Then, PRECiSE-CO96 yields

$$\mathbb{P} \left(\exists t \geq 1 : \max_{\beta \in [-\frac{1}{1-\mu}, \frac{1}{\mu}]} \sum_{i=1}^t \ln(1 + \beta(X_i - \mu)) \geq \ln(\frac{1}{\delta}) + \frac{1}{2} \ln(4t) \right) \leq \delta.$$

Compare with KT guarantee: $\mathbb{P} \left(\exists t \geq 1 : \frac{(\sum_{i=1}^t (X_i - \mu))^2}{2t} \geq \ln(\frac{1}{\delta}) + \frac{1}{2} \ln(4t) \right) \leq \delta.$

✗ No closed-form concentration, need numerical inversion through bisection

✓ **Never-vacuous** for any sample size

PRECiSE-CO96: why never-vacuous?

PRECiSE-CO96: why never-vacuous?

Prior placed on portfolio weight $b \in [0, 1]$ is $\text{Beta}(\frac{1}{2}, \frac{1}{2})$: $F(b) = \frac{1}{\pi\sqrt{b(1-b)}}$.

This has several advantages:

PRECiSE-CO96: why never-vacuous?

Prior placed on portfolio weight $b \in [0, 1]$ is $\text{Beta}(\frac{1}{2}, \frac{1}{2})$: $F(b) = \frac{1}{\pi\sqrt{b(1-b)}}$.

This has several advantages:

- Maps to **full admissible betting range** $\beta = -\frac{1}{1-m} + \left(\frac{1}{1-m} + \frac{1}{m}\right)b \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$ which depends on m

PRECiSE-CO96: why never-vacuous?

Prior placed on portfolio weight $b \in [0, 1]$ is $\text{Beta}(\frac{1}{2}, \frac{1}{2})$: $F(b) = \frac{1}{\pi\sqrt{b(1-b)}}$.

This has several advantages:

- Maps to **full admissible betting range** $\beta = -\frac{1}{1-m} + \left(\frac{1}{1-m} + \frac{1}{m}\right)b \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$ which depends on m
- Allocates **sufficient mass near extreme bets** $\Rightarrow$ prevents early-time regret explosion $\Rightarrow$ CS non-vacuous for all t

PRECiSE-CO96: why never-vacuous?

Prior placed on portfolio weight $b \in [0, 1]$ is $\text{Beta}(\frac{1}{2}, \frac{1}{2})$: $F(b) = \frac{1}{\pi\sqrt{b(1-b)}}$.

This has several advantages:

- Maps to **full admissible betting range** $\beta = -\frac{1}{1-m} + \left(\frac{1}{1-m} + \frac{1}{m}\right)b \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]$ which depends on m
- Allocates **sufficient mass near extreme bets** $\Rightarrow$ prevents early-time regret explosion $\Rightarrow$ CS non-vacuous for all t
- Deeper connections to universal coding, Jefferys prior

PRECiSE-CO96 yields time-uniform empirical Bernstein bound

PRECiSE-CO96 yields time-uniform empirical Bernstein bound

The implicit concentration in Theorem 1 implies an **empirical-Bernstein-type** CS, in contrast to the **Hoeffding-type** CS based on KT.

PRECiSE-CO96 yields time-uniform empirical Bernstein bound

The implicit concentration in Theorem 1 implies an **empirical-Bernstein-type** CS, in contrast to the **Hoeffding-type** CS based on KT. In detail:

Theorem 2. Under the assumptions of Theorem 1, denote by $\hat{\mu}_i = \frac{1}{i} \sum_{j=1}^i X_j$, $V_i = \sum_{j=1}^i (X_j - \hat{\mu}_i)^2$. Then,

$$\mathbb{P}\left(\mu \in \left[\max_{1 \leq i \leq t} \hat{\mu}_i - \varepsilon_i, \min_{1 \leq i \leq t} \hat{\mu}_i + \varepsilon_i\right], \forall t \geq 1\right) \geq 1 - \delta,$$

where for large t , $\varepsilon_t \approx \frac{1}{t} \sqrt{2V_t(\frac{1}{2} \ln t + \ln \frac{1}{\delta})} + \frac{4}{3t}(\frac{1}{2} \ln t + \ln \frac{1}{\delta})$.

PRECiSE-CO96 yields time-uniform empirical Bernstein bound

The implicit concentration in Theorem 1 implies an **empirical-Bernstein-type** CS, in contrast to the **Hoeffding-type** CS based on KT. In detail:

Theorem 2. Under the assumptions of Theorem 1, denote by $\hat{\mu}_i = \frac{1}{i} \sum_{j=1}^i X_j$, $V_i = \sum_{j=1}^i (X_j - \hat{\mu}_i)^2$. Then,

$$\mathbb{P}\left(\mu \in \left[\max_{1 \leq i \leq t} \hat{\mu}_i - \varepsilon_i, \min_{1 \leq i \leq t} \hat{\mu}_i + \varepsilon_i\right], \forall t \geq 1\right) \geq 1 - \delta,$$

where for large t , $\varepsilon_t \approx \frac{1}{t} \sqrt{2V_t(\frac{1}{2} \ln t + \ln \frac{1}{\delta})} + \frac{4}{3t}(\frac{1}{2} \ln t + \ln \frac{1}{\delta})$.

Remarks:

PRECiSE-CO96 yields time-uniform empirical Bernstein bound

The implicit concentration in Theorem 1 implies an **empirical-Bernstein-type** CS, in contrast to the **Hoeffding-type** CS based on KT. In detail:

Theorem 2. Under the assumptions of Theorem 1, denote by $\hat{\mu}_i = \frac{1}{i} \sum_{j=1}^i X_j$, $V_i = \sum_{j=1}^i (X_j - \hat{\mu}_i)^2$. Then,

$$\mathbb{P}\left(\mu \in \left[\max_{1 \leq i \leq t} \hat{\mu}_i - \varepsilon_i, \min_{1 \leq i \leq t} \hat{\mu}_i + \varepsilon_i\right], \forall t \geq 1\right) \geq 1 - \delta,$$

where for large t , $\varepsilon_t \approx \frac{1}{t} \sqrt{2V_t(\frac{1}{2} \ln t + \ln \frac{1}{\delta})} + \frac{4}{3t}(\frac{1}{2} \ln t + \ln \frac{1}{\delta})$.

Remarks:

- Note $\frac{1}{t} \sqrt{V_t \ln \frac{1}{\delta}}$ is the empirical Bernstein rate for fixed t , and $\frac{1}{t} \sqrt{V_t \ln t}$ is price for time-uniformity

PRECiSE-CO96 yields time-uniform empirical Bernstein bound

The implicit concentration in Theorem 1 implies an **empirical-Bernstein-type** CS, in contrast to the **Hoeffding-type** CS based on KT. In detail:

Theorem 2. Under the assumptions of Theorem 1, denote by $\hat{\mu}_i = \frac{1}{i} \sum_{j=1}^i X_j$, $V_i = \sum_{j=1}^i (X_j - \hat{\mu}_i)^2$. Then,

$$\mathbb{P}\left(\mu \in \left[\max_{1 \leq i \leq t} \hat{\mu}_i - \varepsilon_i, \min_{1 \leq i \leq t} \hat{\mu}_i + \varepsilon_i\right], \forall t \geq 1\right) \geq 1 - \delta,$$

where for large t , $\varepsilon_t \approx \frac{1}{t} \sqrt{2V_t(\frac{1}{2} \ln t + \ln \frac{1}{\delta})} + \frac{4}{3t}(\frac{1}{2} \ln t + \ln \frac{1}{\delta})$.

Remarks:

- Note $\frac{1}{t} \sqrt{V_t \ln \frac{1}{\delta}}$ is the empirical Bernstein rate for fixed t , and $\frac{1}{t} \sqrt{V_t \ln t}$ is price for time-uniformity
- **Next:** Going from $\frac{1}{t} \sqrt{V_t \ln t}$ (i.e., CLT rate $\times \sqrt{\ln t}$) to $\frac{1}{t} \sqrt{V_t \ln \ln t}$ (i.e., law of the iterated logarithm (LIL) rate, asymptotically optimal)

Outline

1. Recap and warm-up

- 1.1 Testing by betting
- 1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

- 2.1 A reformulation using online learning tools
- 2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

From PRECiSE-CO96 to -R70: motivation for a different prior

From PRECiSE-CO96 to -R70: motivation for a different prior

By the first-order condition, we have

$$\beta_t^*(m) \in \arg \max_{\beta \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]} \sum_{i=1}^t \ln(1 + \beta(X_i - m)) \approx \frac{\sum_{i=1}^t (X_i - m)}{\sum_{i=1}^t (X_i - m)^2}.$$

From PRECiSE-CO96 to -R70: motivation for a different prior

By the first-order condition, we have

$$\beta_t^*(m) \in \arg \max_{\beta \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]} \sum_{i=1}^t \ln(1 + \beta(X_i - m)) \approx \frac{\sum_{i=1}^t (X_i - m)}{\sum_{i=1}^t (X_i - m)^2}.$$

Under $m = \mu$, the LIL gives $\sum_{i=1}^t (X_i - \mu) = \mathcal{O}(\sqrt{V_t \ln \ln t})$, which implies

$$\beta_t^*(\mu) = \mathcal{O}\left(\sqrt{\frac{\ln \ln t}{V_t}}\right) \rightarrow 0.$$

From PRECiSE-CO96 to -R70: motivation for a different prior

By the first-order condition, we have

$$\beta_t^*(m) \in \arg \max_{\beta \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]} \sum_{i=1}^t \ln(1 + \beta(X_i - m)) \approx \frac{\sum_{i=1}^t (X_i - m)}{\sum_{i=1}^t (X_i - m)^2}.$$

Under $m = \mu$, the LIL gives $\sum_{i=1}^t (X_i - \mu) = \mathcal{O}(\sqrt{V_t \ln \ln t})$, which implies

$$\beta_t^*(\mu) = \mathcal{O}\left(\sqrt{\frac{\ln \ln t}{V_t}}\right) \rightarrow 0.$$

Take-away message: The optimal bet $\rightarrow 0$ as $t \rightarrow \infty$, so a prior must maintain mass near 0 at all times; one that places mass away from 0 pays extra regret.

From PRECiSE-CO96 to -R70: motivation for a different prior

By the first-order condition, we have

$$\beta_t^*(m) \in \arg \max_{\beta \in \left[-\frac{1}{1-m}, \frac{1}{m}\right]} \sum_{i=1}^t \ln(1 + \beta(X_i - m)) \approx \frac{\sum_{i=1}^t (X_i - m)}{\sum_{i=1}^t (X_i - m)^2}.$$

Under $m = \mu$, the LIL gives $\sum_{i=1}^t (X_i - \mu) = \mathcal{O}(\sqrt{V_t \ln \ln t})$, which implies

$$\beta_t^*(\mu) = \mathcal{O}\left(\sqrt{\frac{\ln \ln t}{V_t}}\right) \rightarrow 0.$$

Take-away message: The **optimal bet** $\rightarrow 0$ as $t \rightarrow \infty$, so a prior **must maintain mass near 0** at all times; one that places mass away from 0 pays extra regret.

$\Rightarrow$ Instead of $F(b) = \text{Beta}(\frac{1}{2}, \frac{1}{2})$, PRECiSE-R70 uses the following $F(\beta)$ from Robbins (1970)

$$F(\beta) \asymp \frac{1}{|\beta| \ln\left(\frac{1}{|\beta|}\right) \left(\ln \ln\left(\frac{1}{|\beta|}\right)\right)^2}.$$

Note $F(\beta)$ is a **proper prior with** $F(\beta) \rightarrow \infty$ as $\beta \rightarrow 0$.

PRECiSE-R70:

Portfolio REgret for Confidence SEquences Using Robbins (1970)

PRECiSE-R70:

Portfolio REgret for Confidence SEquences Using Robbins (1970)

Theorem 3. Let $\delta \in (0, 1)$. Let $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_i \mid X_{1:i-1}] = \mu$ a.s.. Denote by $\hat{\mu}_t = \frac{1}{t} \sum_{i=1}^t X_i$, $V_t = \sum_{i=1}^t (X_i - \hat{\mu}_t)^2$. Then, PRECiSE-R70 yields a CS such that

$$\mathbb{P}(\mu \in [\ell_t, u_t], \forall t \geq 1) \geq 1 - \delta,$$

where the corresponding confidence width satisfies

$$\max\{u_t - \hat{\mu}_t, \hat{\mu}_t - \ell_t\} \lesssim \frac{1}{t} \sqrt{2V_t(\ln \ln t + \ln \frac{1}{\delta})}.$$

PRECiSE-R70:

Portfolio REgret for Confidence SEquences Using Robbins (1970)

Theorem 3. Let $\delta \in (0, 1)$. Let $X_1, X_2, \dots \in [0, 1]$ such that $\mathbb{E}[X_i \mid X_{1:i-1}] = \mu$ a.s.. Denote by $\hat{\mu}_t = \frac{1}{t} \sum_{i=1}^t X_i$, $V_t = \sum_{i=1}^t (X_i - \hat{\mu}_t)^2$. Then, PRECiSE-R70 yields a CS such that

$$\mathbb{P}(\mu \in [\ell_t, u_t], \forall t \geq 1) \geq 1 - \delta,$$

where the corresponding confidence width satisfies

$$\max\{u_t - \hat{\mu}_t, \hat{\mu}_t - \ell_t\} \lesssim \frac{1}{t} \sqrt{2V_t(\ln \ln t + \ln \frac{1}{\delta})}.$$

Compare with

- Empirical Bernstein: $|\hat{\mu}_t - \mu| \lesssim \frac{1}{t} \sqrt{2V_t \ln \frac{1}{\delta}}$ for fixed t
- PRECiSE-CO96: $|\hat{\mu}_t - \mu| \lesssim \frac{1}{t} \sqrt{2V_t(\ln t + \ln \frac{1}{\delta})}$ uniformly in t

Outline

1. Recap and warm-up

- 1.1 Testing by betting
- 1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

- 2.1 A reformulation using online learning tools
- 2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

Simulations





Figure: Time-uniform CSs (left) and their widths (right) for a sequence of **i.i.d. Bernoulli(0.1)** r.v.s.

Simulations





Figure: Time-uniform CSs (left) and their widths (right) for a sequence of **i.i.d. Bernoulli(0.1)** r.v.s.

Simulations





Figure: Time-uniform CSs (left) and their widths (right) for a sequence of **i.i.d. Bernoulli(0.5)** r.v.s.

Simulations





Figure: Time-uniform CSs (left) and their widths (right) for a sequence of i.i.d. $\text{Beta}(1,1)$ r.v.s.

Outline

1. Recap and warm-up

- 1.1 Testing by betting
- 1.2 Warm-up with the Krichevsky-Trofimov (KT) betting algorithm

2. Reduction from betting to online portfolio selection

- 2.1 A reformulation using online learning tools
- 2.2 From betting to online portfolio selection

3. PRECiSE-CO96: never-vacuous CS for any sample size T

4. PRECiSE-R70: CS with log-log rate asymptotically in T

5. Simulations

6. Summary

Summary

Summary

1. Betting $\rightarrow$ Ville $\rightarrow$ CS is mechanically simple once you have an e-process $\text{Wealth}_t(m)$

Summary

1. Betting $\rightarrow$ Ville $\rightarrow$ CS is mechanically simple once you have an e-process $\text{Wealth}_t(m)$
2. The paper's novelty: center the concentration inequality at log-wealth **at best bet in hindsight** and control the gap by **universal portfolio regret**

Summary

1. Betting $\rightarrow$ Ville $\rightarrow$ CS is mechanically simple once you have an e-process $\text{Wealth}_t(m)$
2. The paper's novelty: center the concentration inequality at log-wealth **at best bet in hindsight** and control the gap by **universal portfolio regret**
3. We do **not** need to run the ideal betting strategy: a regret bound implies the existence of a CS and enables inversion

Summary

1. Betting $\rightarrow$ Ville $\rightarrow$ CS is mechanically simple once you have an e-process $\text{Wealth}_t(m)$
2. The paper's novelty: center the concentration inequality at log-wealth **at best bet in hindsight** and control the gap by **universal portfolio regret**
3. We do **not** need to run the ideal betting strategy: a regret bound implies the existence of a CS and enables inversion
4. **PRECiSE-CO96** uses $\text{Beta}(\frac{1}{2}, \frac{1}{2})$ prior on the portfolio weight $b \in [0, 1]$
 $\Rightarrow$ first **never-vacuous** CS for *any* sample size T

Summary

1. Betting $\rightarrow$ Ville $\rightarrow$ CS is mechanically simple once you have an e-process $\text{Wealth}_t(m)$
2. The paper's novelty: center the concentration inequality at log-wealth **at best bet in hindsight** and control the gap by **universal portfolio regret**
3. We do **not** need to run the ideal betting strategy: a regret bound implies the existence of a CS and enables inversion
4. **PRECiSE-CO96** uses $\text{Beta}(\frac{1}{2}, \frac{1}{2})$ prior on the portfolio weight $b \in [0, 1]$
 $\Rightarrow$ first **never-vacuous** CS for *any* sample size T
5. **PRECiSE-R70** uses a **zero-centered prior** on the bet $\beta \in [-1, 1]$
 $\Rightarrow$ LIL rate asymptotically

References

- Cover, T. and Ordentlich, E. (1996). Universal portfolios with side information. *IEEE Transactions on Information Theory*, 42(2):348–363.
- Cover, T. M. (1991). Universal portfolios. *Mathematical Finance*, 1(1):1–29.
- Orabona, F. and Jun, K.-S. (2024). Tight concentrations and confidence sequences from the regret of universal portfolio. *IEEE Transactions on Information Theory*, 70(1):436–455.
- Orabona, F. and Pal, D. (2016). Coin Betting and Parameter-Free Online Learning. In Lee, D., Sugiyama, M., Luxburg, U., Guyon, I., and Garnett, R., editors, *Advances in Neural Information Processing Systems*, volume 29. Curran Associates, Inc.
- Ramdas, A. and Wang, R. (2025). Hypothesis testing with e-values. *Foundations and Trends® in Statistics*, 1(1-2):1–390.
- Robbins, H. (1970). Statistical Methods Related to the Law of the Iterated Logarithm. *The Annals of Mathematical Statistics*, 41(5):1397 – 1409.